

Highlights of Linear Algebra

Vectors, Matrices, and Geometric Intuition

Firuz Kamalov

MAI 600: Mathematics and Statistics for AI

September 12, 2025

Why is Linear Algebra the Bedrock of AI?

1. It's the Language of Data

Linear algebra gives us a universal way to represent and work with data.

- **Vectors** represent individual data points (a user, a word, a patient).
- **Matrices** represent entire datasets or systems (a collection of users, an image, the weights of a neural network).

Why is Linear Algebra the Bedrock of AI?

1. It's the Language of Data

Linear algebra gives us a universal way to represent and work with data.

- **Vectors** represent individual data points (a user, a word, a patient).
- **Matrices** represent entire datasets or systems (a collection of users, an image, the weights of a neural network).

2. It's the Engine of Computation

Many AI algorithms are fundamentally a series of linear algebra operations.

- **Deep Learning:** A neural network's forward pass is just a sequence of matrix multiplications and vector additions.
- **Dimensionality Reduction (PCA):** Finds the most important "directions" in the data.

Key AI Applications

- **Computer Vision:** Images are matrices of pixels. Rotations, scaling, and filtering are all matrix operations (<https://www.youtube.com/watch?v=aircAruvnKk>).
- **Natural Language Processing (NLP):** Word embeddings like Word2Vec represent words as vectors, capturing semantic meaning.
- **Recommender Systems:** Analyze user-item matrices to predict what you might like next (e.g., Netflix, Amazon).
<https://www.youtube.com/watch?v=aircAruvnKk>

- 1 Vectors: The Data Atoms
- 2 Vector Spaces, Norms, and Inner Products
- 3 Matrices and Linear Transformations

What is a Vector?

A vector is an object with both **magnitude** (length) and **direction**. We can think of it in two ways:


Algebraically

An ordered list of numbers. In AI, we typically write vectors as **column vectors**.

$$\mathbf{v} = \begin{bmatrix} 3 \\ 2 \end{bmatrix} \quad \mathbf{w} = \begin{bmatrix} 0.8 \\ 1.2 \\ -3.5 \\ 0.1 \end{bmatrix}$$

Geometrically

An arrow in space pointing from the origin to a coordinate point.




Example: A Feature Vector

A vector can represent a single data point. For a housing dataset:

$$\text{house} = \begin{bmatrix} \text{Area (sq ft)} \\ \text{Bedrooms} \\ \text{Age (years)} \end{bmatrix} = \begin{bmatrix} 1500 \\ 3 \\ 10 \end{bmatrix}$$


Vector Addition (Tip-to-Tail Rule)

$$\begin{bmatrix} a \\ b \end{bmatrix} + \begin{bmatrix} c \\ d \end{bmatrix} = \begin{bmatrix} a + c \\ b + d \end{bmatrix}$$



Scalar Multiplication (Stretching/Shrinking)

$$c \begin{bmatrix} a \\ b \end{bmatrix} = \begin{bmatrix} ca \\ cb \end{bmatrix}$$



Problem

Given the vectors $\mathbf{v} = \begin{bmatrix} 2 \\ 3 \end{bmatrix}$ and $\mathbf{w} = \begin{bmatrix} 3 \\ -1 \end{bmatrix}$.


- 1 Calculate the resultant vector $\mathbf{z} = \mathbf{v} + \mathbf{w}$.
- 2 On a 2D grid, draw $\mathbf{v}$, $\mathbf{w}$, and $\mathbf{z}$. Show the tip-to-tail addition graphically.

Solution

1. Calculation:

$$z = \begin{bmatrix} 2 \\ 3 \end{bmatrix} + \begin{bmatrix} 3 \\ -1 \end{bmatrix} = \begin{bmatrix} 2+3 \\ 3-1 \end{bmatrix} = \begin{bmatrix} 5 \\ 2 \end{bmatrix}$$

2. Drawing:



Definition: A set V closed under addition and scalar multiplication.

Example: $\mathbb{R}^2$ with vectors (x, y) .

- $u = (1, 2), v = (3, 4)$ - $u + v = (4, 6)$ - $2u = (2, 4)$

This stays inside $\mathbb{R}^2$, so it's a vector space.

Span: All linear combinations of vectors.

Example: $\text{Span}\{(1, 0), (0, 1)\} = \mathbb{R}^2$.

AI Link: Features in ML can be represented as spanning vectors in a high-dimensional space.

Example: Word Embedding Space

In NLP, we don't have just 2 basis vectors, but perhaps 300! Each basis vector might represent some abstract linguistic feature. A word like "learning" is then defined by its unique combination of these 300 basis vectors. The entire vocabulary lives within this 300-dimensional vector space.

Vector Norms: Measuring Length

A **norm** is a function that measures the length or magnitude of a vector. It's a key tool for measuring error or distance.

L2 Norm (Euclidean Distance)

The most common "straight-line" distance.

$$\|\mathbf{v}\|_2 = \sqrt{v_1^2 + v_2^2 + \cdots + v_n^2}$$

Used everywhere, from calculating the error of a prediction to finding nearest neighbors.

L1 Norm (Manhattan Distance)

The "city block" distance; sum of absolute components.

$$\|\mathbf{v}\|_1 = |v_1| + |v_2| + \cdots + |v_n|$$

Often used in machine learning when we want to encourage sparsity (many zero values).

Vector Norms: Measuring Length

A **norm** is a function that measures the length or magnitude of a vector. It's a key tool for measuring error or distance.

L2 Norm (Euclidean Distance)

The most common "straight-line" distance.

$$\|\mathbf{v}\|_2 = \sqrt{v_1^2 + v_2^2 + \cdots + v_n^2}$$


Used everywhere, from calculating the error of a prediction to finding nearest neighbors.

L1 Norm (Manhattan Distance)

The "city block" distance; sum of absolute components.

$$\|\mathbf{v}\|_1 = |v_1| + |v_2| + \cdots + |v_n|$$

Often used in machine learning when we want to encourage sparsity (many zero values).



Problem

Imagine a simple feature vector for a movie recommendation system is $\mathbf{v} = \begin{bmatrix} 2 \\ -3 \\ 6 \end{bmatrix}$, representing scores for action, comedy, and drama.

- 1 Calculate the L2 norm (Euclidean length) of the vector, $\|\mathbf{v}\|_2$.
- 2 Calculate the L1 norm (Manhattan length) of the vector, $\|\mathbf{v}\|_1$.

Solution

Given $\mathbf{v} = \begin{bmatrix} 2 \\ -3 \\ 6 \end{bmatrix}$:

Solution

Given $\mathbf{v} = \begin{bmatrix} 2 \\ -3 \\ 6 \end{bmatrix}$:

① **L2 Norm:**

$$\|\mathbf{v}\|_2 = \sqrt{2^2 + (-3)^2 + 6^2} = 7$$

Solution

Given $\mathbf{v} = \begin{bmatrix} 2 \\ -3 \\ 6 \end{bmatrix}$:

① **L2 Norm:**

$$\|\mathbf{v}\|_2 = \sqrt{2^2 + (-3)^2 + 6^2} = 7$$

② **L1 Norm:**

$$\|\mathbf{v}\|_1 = |2| + |-3| + |6| = 11$$

The Inner Product: Projection and Similarity

The **inner product** (or **dot product**) takes two vectors and returns a single scalar.

$$\mathbf{v} \cdot \mathbf{w} = \mathbf{v}^T \mathbf{w} = v_1 w_1 + v_2 w_2 + \cdots + v_n w_n$$

Example


Let $\mathbf{v} = \begin{bmatrix} 2 \\ 3 \\ -1 \end{bmatrix}$ and $\mathbf{w} = \begin{bmatrix} 5 \\ 4 \\ 2 \end{bmatrix}$. Then the dot product is $\mathbf{v} \cdot \mathbf{w} = 10 + 12 + (-2) = 20$.

Its real power comes from its geometric interpretation:


$$\mathbf{v} \cdot \mathbf{w} = \|\mathbf{v}\|_2 \|\mathbf{w}\|_2 \cos(\theta)$$

The inner product (or dot product) $\mathbf{v} \cdot \mathbf{w}$ tells us about the **alignment** of two vectors. Geometrically, it's related to how much of one vector "points in the direction of" or "projects onto" the other.

Inner Product: Geometric Intuition (Projection)



Inner Product: Geometric Intuition (Projection)



Key Takeaway

The quantity $\|\mathbf{v}\| \cos(\theta)$ is the **length** of the projection of $\mathbf{v}$ onto $\mathbf{w}$. The dot product is simply this length multiplied by the length of $\mathbf{w}$.

- If $\mathbf{v}$ and $\mathbf{w}$ point in the same direction ($\theta = 0^\circ$), $\cos(\theta) = 1$, dot product is maximal.
- If they are perpendicular ($\theta = 90^\circ$), $\cos(\theta) = 0$, dot product is zero.

Application: Cosine Similarity

By rearranging the formula, we get a direct measure of similarity that ignores vector magnitudes:

$$\cos(\theta) = \frac{\mathbf{v} \cdot \mathbf{w}}{\|\mathbf{v}\|_2 \|\mathbf{w}\|_2}$$

This is the engine behind recommender systems ("users who liked X also liked Y") and document search. It measures the similarity of direction of user preference vectors or document topic vectors.

Problem

Let's model two users' movie ratings. User A's ratings for (Star Wars, Titanic) are $\mathbf{a} = \begin{bmatrix} 5 \\ 1 \end{bmatrix}$.

User B's ratings are $\mathbf{b} = \begin{bmatrix} 4 \\ 2 \end{bmatrix}$.

- 1 Calculate the inner product $\mathbf{a} \cdot \mathbf{b}$.
- 2 Calculate the cosine similarity between $\mathbf{a}$ and $\mathbf{b}$. Do these users have similar taste?

Solution

① Inner Product:

$$\mathbf{a} \cdot \mathbf{b} = (5)(4) + (1)(2) = 22$$

Since the result is a large positive number, we know the vectors are generally pointing in the same direction.

Solution

① Inner Product:

$$\mathbf{a} \cdot \mathbf{b} = (5)(4) + (1)(2) = 22$$

Since the result is a large positive number, we know the vectors are generally pointing in the same direction.

② Cosine Similarity: First, we need the norms.

- $\|\mathbf{a}\|_2 = \sqrt{5^2 + 1^2} = \sqrt{26}$
- $\|\mathbf{b}\|_2 = \sqrt{4^2 + 2^2} = \sqrt{20}$

Solution

① Inner Product:

$$\mathbf{a} \cdot \mathbf{b} = (5)(4) + (1)(2) = 22$$

Since the result is a large positive number, we know the vectors are generally pointing in the same direction.

② Cosine Similarity: First, we need the norms.

- $\|\mathbf{a}\|_2 = \sqrt{5^2 + 1^2} = \sqrt{26}$
- $\|\mathbf{b}\|_2 = \sqrt{4^2 + 2^2} = \sqrt{20}$

Now, apply the formula:

$$\cos(\theta) = \frac{\mathbf{a} \cdot \mathbf{b}}{\|\mathbf{a}\| \|\mathbf{b}\|} = \frac{22}{\sqrt{26}\sqrt{20}} \approx 0.96$$

A cosine similarity of 0.96 is very close to 1, indicating that the users have very similar taste profiles.

A matrix is a rectangular grid of numbers. In AI, they have two primary roles:

1. As containers for data

- An entire dataset, where rows are samples and columns are features.
- A grayscale image, where each element is a pixel intensity.
- A batch of data fed into a neural network.

$$\text{Data} = \begin{bmatrix} \text{— house 1 —} \\ \text{— house 2 —} \\ \vdots \\ \text{— house m —} \end{bmatrix}$$

2. As linear transformations

A matrix can be seen as a function that transforms vectors from one space to another.

- The weights of a neural network layer form a matrix that transforms the input data.
- Matrices can rotate, scale, and shear vectors.

$$\mathbf{y} = A\mathbf{x} \quad (\text{Matrix } A \text{ "acts on" vector } \mathbf{x})$$

Addition and Scalar Multiplication

Just like vectors, matrices can be added (if they have the same dimensions) and multiplied by scalars. This is done element-wise.

$$\begin{bmatrix} a & b \\ c & d \end{bmatrix} + \begin{bmatrix} e & f \\ g & h \end{bmatrix} = \begin{bmatrix} a+e & b+f \\ c+g & d+h \end{bmatrix} \quad k \begin{bmatrix} a & b \\ c & d \end{bmatrix} = \begin{bmatrix} ka & kb \\ kc & kd \end{bmatrix}$$

Problem

Given two matrices A and B , which might represent two different sets of features for the same batch of data:

$$A = \begin{bmatrix} 1 & 5 \\ 3 & 0 \\ -2 & 4 \end{bmatrix} \quad B = \begin{bmatrix} 0 & 2 \\ -1 & 7 \\ 5 & -3 \end{bmatrix}$$

Calculate the sum $A + B$.

Solution

Matrix addition is performed element-wise:

$$\begin{aligned} A + B &= \begin{bmatrix} 1 & 5 \\ 3 & 0 \\ -2 & 4 \end{bmatrix} + \begin{bmatrix} 0 & 2 \\ -1 & 7 \\ 5 & -3 \end{bmatrix} \\ &= \begin{bmatrix} 1+0 & 5+2 \\ 3+(-1) & 0+7 \\ -2+5 & 4+(-3) \end{bmatrix} \\ &= \begin{bmatrix} 1 & 7 \\ 2 & 7 \\ 3 & 1 \end{bmatrix} \end{aligned}$$

The resulting matrix has the same dimensions as A and B .

The Most Important Operation: Matrix Multiplication

Matrix multiplication is how transformations are applied and combined.

Matrix-Vector Multiplication: Transforms a vector.

$$\begin{bmatrix} a & b \\ c & d \end{bmatrix} \begin{bmatrix} x \\ y \end{bmatrix} = \begin{bmatrix} ax + by \\ cx + dy \end{bmatrix}$$

The result is a linear combination of the columns of the matrix.

Matrix-Matrix Multiplication: Represents a composition of transformations. If $C = AB$, applying matrix C is the same as applying B first, then A .

Example: Matrix-Vector Multiplication

Problem

Suppose a neural network layer has a weight matrix W and receives an input vector $\mathbf{x}$:

$$W = \begin{bmatrix} 2 & -1 & 0 \\ 1 & 3 & -2 \end{bmatrix} \quad \mathbf{x} = \begin{bmatrix} 4 \\ 1 \\ 5 \end{bmatrix}$$

Calculate the output vector $\mathbf{y} = W\mathbf{x}$ before the bias and activation function are applied.

Solution

We multiply each row of the matrix by the column vector:

$$\begin{aligned}\mathbf{y} = W\mathbf{x} &= \begin{bmatrix} 2 & -1 & 0 \\ 1 & 3 & -2 \end{bmatrix} \begin{bmatrix} 4 \\ 1 \\ 5 \end{bmatrix} \\ &= \begin{bmatrix} (2)(4) + (-1)(1) + (0)(5) \\ (1)(4) + (3)(1) + (-2)(5) \end{bmatrix} \\ &= \begin{bmatrix} 7 \\ -3 \end{bmatrix}\end{aligned}$$

The input vector in $\mathbb{R}^3$ was transformed into an output vector in $\mathbb{R}^2$.

Example: Matrix-Matrix Multiplication

Problem

Given two matrices A and B :

$$A = \begin{bmatrix} 1 & 2 \\ 3 & 4 \end{bmatrix} \quad B = \begin{bmatrix} 5 & 6 \\ 7 & 8 \end{bmatrix}$$

Calculate the product $C = AB$.

Solution

To calculate $C = AB$, we multiply rows of A by columns of B :

$$\begin{aligned} C &= \begin{bmatrix} 1 & 2 \\ 3 & 4 \end{bmatrix} \begin{bmatrix} 5 & 6 \\ 7 & 8 \end{bmatrix} \\ &= \begin{bmatrix} (1)(5) + (2)(7) & (1)(6) + (2)(8) \\ (3)(5) + (4)(7) & (3)(6) + (4)(8) \end{bmatrix} \\ &= \begin{bmatrix} 19 & 22 \\ 43 & 50 \end{bmatrix} \end{aligned}$$


Remember, matrix multiplication is generally not commutative ($AB \neq BA$).

Geometric Intuition of a Linear Transformation

A matrix transforms the entire vector space. The key insight is to see what it does to the basis vectors $\mathbf{i}$ and $\mathbf{j}$. The columns of a matrix tell you where the basis vectors land after the transformation.

$$A = \begin{bmatrix} 1 & -1 \\ 1 & 1 \end{bmatrix}$$

This matrix sends $\mathbf{i} = \begin{bmatrix} 1 \\ 0 \end{bmatrix}$ to $\begin{bmatrix} 1 \\ 1 \end{bmatrix}$ and $\mathbf{j} = \begin{bmatrix} 0 \\ 1 \end{bmatrix}$ to $\begin{bmatrix} -1 \\ 1 \end{bmatrix}$.



Problem

Consider the "shear" transformation matrix $A = \begin{bmatrix} 1 & 1 \\ 0 & 1 \end{bmatrix}$ and the vector $\mathbf{v} = \begin{bmatrix} 2 \\ 1 \end{bmatrix}$.


- 1 Calculate the transformed vector $\mathbf{w} = A\mathbf{v}$.
- 2 Draw $\mathbf{v}$ and $\mathbf{w}$ on a 2D grid. What is the geometric effect of this transformation?

Solution

1. Calculation:


$$\mathbf{w} = \begin{bmatrix} 1 & 1 \\ 0 & 1 \end{bmatrix} \begin{bmatrix} 2 \\ 1 \end{bmatrix} = \begin{bmatrix} (1)(2) + (1)(1) \\ (0)(2) + (1)(1) \end{bmatrix} = \begin{bmatrix} 3 \\ 1 \end{bmatrix}$$

2. Geometric Effect:



The transformation shifted the top of the vector to the right. It "sheared" the space horizontally.

Shear mapping




Geometric Intuition of a Rotation

A rotation is a linear transformation. We can see what it does by observing its effect on the basis vectors $\mathbf{i}$ and $\mathbf{j}$.

The columns of the rotation matrix tell you where the basis vectors land. For a counter-clockwise rotation by 30° , the matrix is:

$$R_{30} = \begin{bmatrix} \cos(30^\circ) & \cos(120^\circ) \\ \sin(30^\circ) & \sin(120^\circ) \end{bmatrix} = \begin{bmatrix} \frac{\sqrt{3}}{2} & -\frac{1}{2} \\ \frac{1}{2} & \frac{\sqrt{3}}{2} \end{bmatrix}$$

This matrix sends $\mathbf{i} = \begin{bmatrix} 1 \\ 0 \end{bmatrix}$ to $\begin{bmatrix} \sqrt{3}/2 \\ 1/2 \end{bmatrix}$ and $\mathbf{j} = \begin{bmatrix} 0 \\ 1 \end{bmatrix}$ to $\begin{bmatrix} -1/2 \\ \sqrt{3}/2 \end{bmatrix}$.



- **Vectors** are the atoms of our data, representing single data points with magnitude and direction.
- **Matrices** are collections of data or, more powerfully, **linear transformations** that manipulate vectors.
- **Vector Spaces** provide the "universe" where our data lives, and their properties are defined by basis vectors.
- **Norms** measure size/distance, while **Inner Products** measure similarity/projection. These are critical for loss functions and similarity metrics.
- **Matrix Multiplication** is the engine of AI, from applying transformations in computer vision to processing data in neural networks.

Questions?